



Enterprise AI Playground

Design overview

What the Enterprise AI Playground is, how it is built, and how the agents inside it work. Written for people in IT and the business who decide, buy and operate, not only for engineers. Two pages of tables, then the detail you need when you need it.

Governed


Identity, policy and budget before every call; one ledger row after it.

Agentic

Blueprints as explicit steps, review gates for people, evaluation and canaries.

Open

Thirty models, any SDK through one gateway, four cloud runtimes, no lock-in.



Document

Design overview

Date

29 September 2026

Audience

IT and business decision makers




In one screen

What it is	One governed place where an organisation’s people explore every model, run and build agents, and see every token and dollar
Who uses it	Explorers (try and learn), builders (make agents and notebooks), champions (lead a department), administrators (policy, budgets, keys)
Where it runs	Azure Container Apps from one command; the same images run on any container platform; sign-in through Entra, Google, Okta or any OpenID Connect provider
What it costs to run	About 90 to 150 USD a month at pilot sizing, plus model usage, which the ledger shows per person
What makes it enterprise-grade	Every call passes identity, policy and budget before it runs; every run can pause for a person; every step is traced; every agent is evaluated nightly and rolled back on drift



How a request travels




How a request travels through the layers. The animated version plays in the product under Documentation, Design overview.

LAYER	WHAT IT DOES	WHAT YOU GET
People and clients	The web app, OpenAI-compatible SDKs, MCP clients (Claude, Cursor, VS Code) and low-code studios all enter through the same door	One set of rules whatever tool a person prefers

4. **Reason.** The blueprint runs its steps: deterministic code wherever the answer can be computed, the model only where judgement is needed, retrieval with citations where facts are needed.
5. **Review.** On an exception the run pauses at a human gate; a reviewer approves or rejects and the run resumes from its checkpoint.
6. **Answer.** The output comes with its trace: every step, token, dollar and citation, visible in Runs, Traces and Cost.
7. **Improve.** Golden cases graded by a judge model gate the agent's hardening level; the nightly canary demotes and rolls back on drift.
8. **Share.** The same run becomes a notebook, a framework project for any SDK, a low-code flow, or a step-by-step cloud deployment.

The agents, and what happens behind the scenes

Every agent is a blueprint: an explicit graph of steps with three kinds of node. **Tool** steps are ordinary code (load, extract, compute, retrieve) and never call a model. **Model** steps call a model with a narrow brief. **Gate** steps decide, and **human** gates pause. This split is what keeps cost, latency and behaviour predictable: the model does only the parts that need language or judgement.



Inside an agent: code steps, a rule, a person on exceptions and one model call.

AGENT	WHAT IT IS FOR	BEHIND THE SCENES	WHERE A PERSON DECIDES
Document reconciliation	Match supplier invoices to purchase orders and contracts before payment	Load documents → extract fields → validate against tolerances → quality gate → write a summary of exceptions	The human decision step on any invoice that fails validation

AGENT	WHAT IT IS FOR	BEHIND THE SCENES	WHERE A PERSON DECIDES
Sage Lens deep research	Sourced research on a question	Clarity gate → research over the corpus → validate coverage → synthesise with sources	A clarification when the question is vague
Learning path generator	Three-level learning paths per role	Generate the path → a critic reviews it → curate references from a vetted list → assemble	Review of the published path
Review panel	Any draft checked before it ships	Load the draft → five critics (legal, consistency, completeness, policy, tone) → human verdict → revise or publish	The verdict
Data analyst	Questions over a CSV with every number computed in code	Load the file → route the question → compute with pandas → narrate the result	None needed: numbers never come from the model
Knowledge Q&A	Grounded answers over policy documents	Retrieve passages → answer with citations → verify every citation exists	Answers without valid citations are flagged

The platform runs itself with the same machinery. **Key health check** sends a tiny completion through every platform key and raises an alert when a provider rejects one. **Cost sentinel** compares yesterday's spend per department and feature with the trailing week and explains anomalies. **Connector reviewer** shortlists pending MCP servers and recommends approve, hold or block. **Onboarding coach** reads what a person has used and suggests their next three steps. **Adoption digest** and **Showcase writer** turn analytics and runs into short posts that a person edits before publishing.

Custom agents built in the wizard use the same runtime: instructions, Knowledge Spaces, skills, built-in tools (calculator, datasets, knowledge search, web fetch, sandboxed Python) and approved MCP connectors, with the same governance and the same traces.

Features, page by page

AREA	PAGE	WHAT IT DOES	WHAT IT OFFERS
Home	Console	Credits, spend, tokens, hours by feature, models used	The state of adoption at a glance

AREA	PAGE	WHAT IT DOES	WHAT IT OFFERS
Home	Documentation	Every guide, page how-to and the API reference, searchable, inside the product	Nothing to look up elsewhere
Discover	Models	30 models across 11 providers with prices, context and capabilities; each runs on a platform key or your own	One price sheet for every model
Discover	Blueprints	175 blueprints across six families, Gen AI or Agentic AI, each with two ways to build	Start from something that already works
Discover	Frameworks	The same blueprint as an OpenAI Agents SDK, LangGraph, CrewAI, Microsoft Agent Framework or Google ADK project, run in the sandbox through the gateway	Leave with code that runs
Discover	Low-code studios	Importable Langflow flows and n8n workflows, Copilot Studio recipes	The no-code route to the same agent
Discover	Cloud platforms	Foundry, AgentCore, Google Agent Runtime and Anthropic Managed Agents with portals, CLI sign-in and step-by-step deploy guides per framework and model	A path to production on any cloud
Discover	MCP Marketplace	7,500 servers with ready-made configuration for seven clients; admins approve before agents may call	Integrations without surprises
Discover	Popular Git repos, Skills	The frameworks, harnesses and reference implementations behind the playground; SKILL.md packs for agents	Learn from the source
Build	Playground	Chat, compare four models side by side, copy the request as code, Smart routing	Fast, governed exploration
Build	Agent Hub	Run blueprints with review gates, build custom agents, take them to any SDK	From idea to working agent
Build	Notebooks	Every blueprint as a notebook that runs in the browser or a sandbox, with GPU options	Reproducible experiments
Build	Knowledge	Knowledge Spaces over documents, pages and repositories with hybrid search and cited answers	Grounded, auditable answers

AREA	PAGE	WHAT IT DOES	WHAT IT OFFERS
Build	Datasets	Eight mock sets with schema and downloads, plus trusted public sources with loaders	Safe data, then real data
Evaluate	Evals, Canary	Golden sets, rubrics, a judge, gates; nightly reruns with drift alerts and rollback	Quality that is measured, not assumed
Operate	Runs, Traces, Cost, Adoption	Every run and its review inbox; every call as a timeline; cost drill-down to the ledger row with CSV export; hours, outcomes and ROI by department	Control without a separate tool
Community	Showcase, Challenges, Leaderboard	Published wins, judged builds, points earned from the ledger	Adoption that spreads on its own
Admin	Users, Policies, Budgets, Settings	Roles, per-role model and tool policy, caps with alerts, platform keys and their scope, personal tokens, MCP client setup	The controls in one place

Security and identity

CONCERN	HOW IT IS HANDLED
Sign-in	Microsoft Entra ID, Google Workspace, Okta or any OpenID Connect provider; e-mail domain and address allow-lists; administrators named by e-mail
Sessions	Stateless, signed; signing out records a time after which no earlier session is accepted
API exposure	The API is reachable only from the web tier with an internal key plus identity headers; never on the public internet
Keys	Platform keys in Key Vault; personal keys encrypted per person; every call records which key paid
Untrusted code	Notebook and framework code runs in isolated sandboxes with egress disabled on Azure
Data	One region; model calls leave only to the providers enabled; MCP servers only after admin approval

CONCERN	HOW IT IS HANDLED
Change control	Every push runs lint, tests, end-to-end checks, secret scanning, dependency audit and template validation before an image exists

Deployment

One script provisions the resource group, registry, images, database, key vault, sessions pool and both container apps on Azure, then waits for the health check; a GitHub Actions workflow builds the images where cloud builds are not permitted. Re-running with a new commit rolls a new revision and keeps the previous one for rollback. Vercel, a single virtual machine and Kubernetes are assessed in the deployment guide; the recommendation is Azure Container Apps for the reasons the guide states.

Glossary

TERM	MEANING
Blueprint	An agent defined as an explicit graph of tool, model, gate and human steps
Review gate	A step where a run pauses until a person approves or rejects
Smart routing	Choosing the cheapest model tier that can do the task, with the saving recorded
Ledger	The table with one row per model call that every cost view reconciles to
Knowledge Space	A searchable, cited collection of documents, pages or repositories
MCP	Model Context Protocol: the standard through which agents call external tools and data
Canary	A nightly re-run of an agent's golden cases to catch drift